

Oct, 2019

# BPF Performance Tools Workshop

Brendan Gregg

*Senior Performance Engineer*

**NETFLIX**



USENIX LISA 2019, Portland, Oct 28-30

# Welcome to BPF Performance Tools

## 1. You can either:

- A) SSH to a lab instance (see bit of paper)
- B) setup your own system. Install BCC & bpftrace, the labs from <https://github.com/brendangregg/bpf-perf-workshop>, and extra tools from <https://github.com/brendangregg/bpf-perf-tools-book>

## 2. Check BCC and bpftrace tools work. On the lab instance:

```
$ sudo bash
# opensnoop-bpcc
[... ]
# opensnoop.bt
[... ]
```

BCC (Ubuntu package naming scheme)

bpftrace

Press Ctrl-C to end each of these

## 3. Begin lab 1. You have 15 minutes.

# Learning Objectives

1. Understand BPF, BCC, and bpftrace
2. Follow different analysis methodologies
3. Use BCC tools to analyze disk I/O issues
4. “ “ short-lived process issues
5. “ “ runq latency issues
6. Develop at least one new bpftrace tool

This was developed as a 90-minute workshop for USENIX LISA 2019

# URLs

- <https://github.com/brendangregg/bpf-perf-workshop>
- <https://github.com/iovisor/bcc> Labs 1-3
  - <https://github.com/iovisor/bcc/blob/master/docs/tutorial.md>
- <https://github.com/iovisor/bpftrace> Labs 4-5
  - [https://github.com/iovisor/bpftrace/blob/master/docs/tutorial\\_one\\_liners.md](https://github.com/iovisor/bpftrace/blob/master/docs/tutorial_one_liners.md)
  - [https://github.com/iovisor/bpftrace/blob/master/docs/reference\\_guide.md](https://github.com/iovisor/bpftrace/blob/master/docs/reference_guide.md)
- <https://github.com/brendangregg/bpf-perf-tools-book>

# Lab 1 Discussion

**BPF**

# BPF 1992: Berkeley Packet Filter


```
# tcpdump -d host 127.0.0.1 and port 80
(000) ldh      [12]
(001) jeq      #0x800          jt 2   jf 18
(002) ld       [26]
(003) jeq      #0x7f000001     jt 6   jf 4
(004) ld       [30]
(005) jeq      #0x7f000001     jt 6   jf 18
(006) ldb      [23]
(007) jeq      #0x84          jt 10  jf 8
(008) jeq      #0x6           jt 10  jf 9
(009) jeq      #0x11          jt 10  jf 18
(010) ldh      [20]
(011) jset     #0x1fff         jt 18  jf 12
(012) ldxb     4*([14]&0xf)
(013) ldh      [x + 14]
(014) jeq      #0x50          jt 17  jf 15
(015) ldh      [x + 16]
(016) jeq      #0x50          jt 17  jf 18
(017) ret      #262144
(018) ret      #0
```

A limited  
**virtual machine** for  
efficient packet filters

# BPF 2019: aka extended BPF




bpfttrace



XDP



BPF microconference



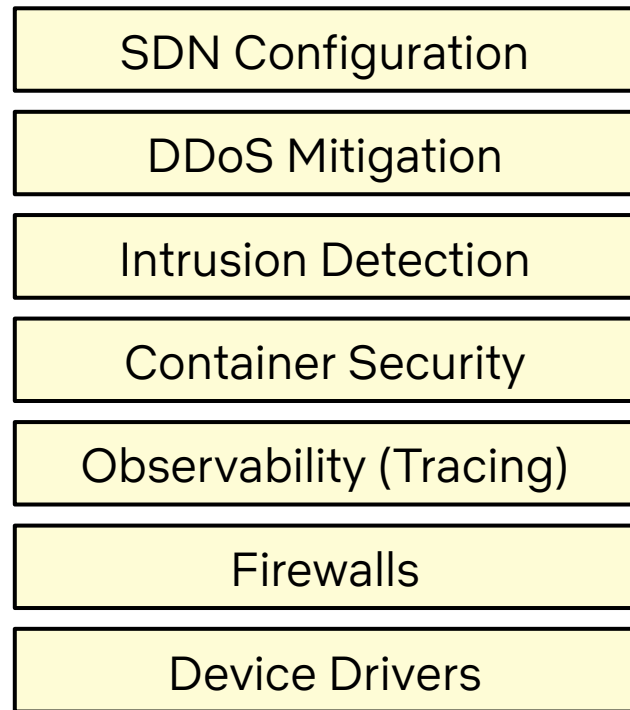
bpfconf

& Facebook Katran, Google KRSI, Netflix flowsrus,  
and many more




# BPF 2019

## User-Defined BPF Programs



...

## Kernel



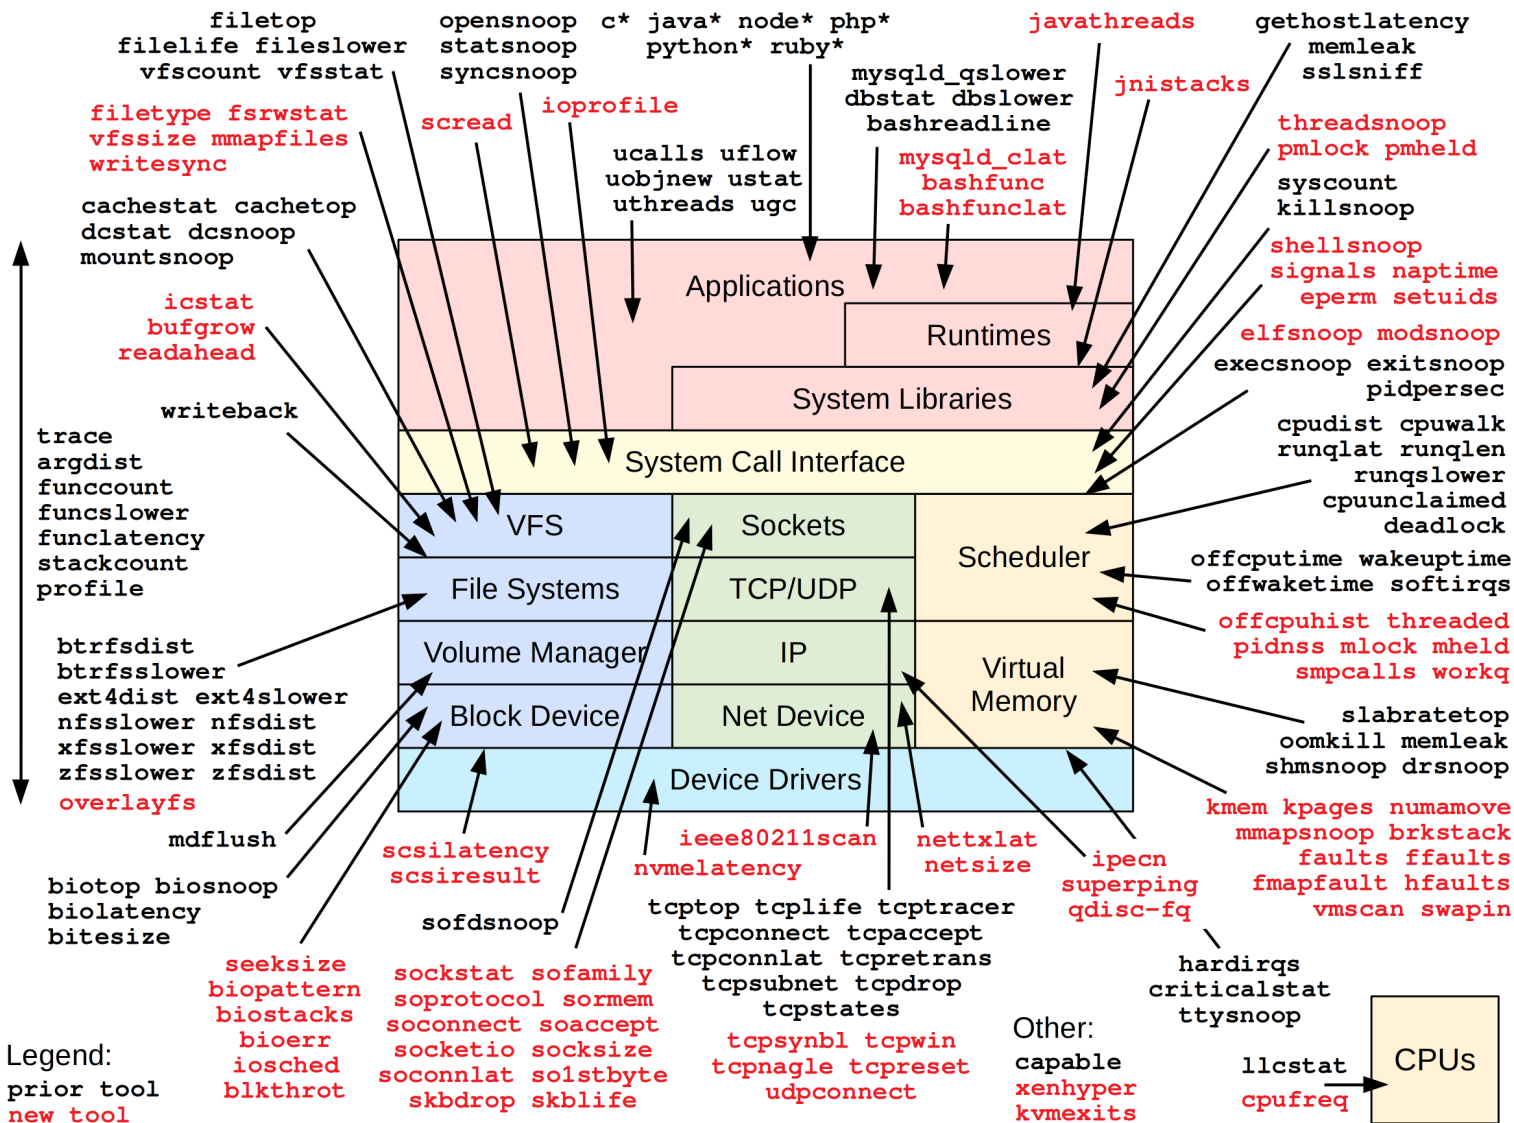

# BPF Recommended Tracing Front Ends

- BCC
  - Great for canned tools & daemons
  - <https://github.com/iovisor/bcc>
- bpftrace
  - Great for developing new short tools & one-liners
  - <https://github.com/iovisor/bpftrace>

The difference is like C vs shell  
Polished products are better in C (BCC)  
Custom one-offs are better in shell (bpftrace)

# BPF Perf Tools (2019)

BCC & bpfftrace repos contain those in black; the book repo has extras in red



# BPF Tool Examples

# CPU: execsnoop

## New process trace

```
# execsnoop.py -T
TIME(s) PCOMM          PID    PPID    RET  ARGS
0.506   run                8745   1828     0   ./run
0.507   bash                8745   1828     0   /bin/bash
0.511   svstat              8747   8746     0   /command/svstat /service/nflx-httpd
0.511   perl                8748   8746     0   /usr/bin/perl -e $1=<>;$1=~/(\\d+) sec;/p...
0.514   ps                  8750   8749     0   /bin/ps --ppid 1 -o pid,cmd,args
0.514   grep                8751   8749     0   /bin/grep org.apache.catalina
0.514   sed                 8752   8749     0   /bin/sed s/^ *//;
0.515   xargs               8754   8749     0   /usr/bin/xargs
0.515   cut                 8753   8749     0   /usr/bin/cut -d  -f 1
0.523   echo                8755   8754     0   /bin/echo
0.524   mkdir               8756   8745     0   /bin/mkdir -v -p /data/tomcat
[...]
1.528   run                8785   1828     0   ./run
1.529   bash                8785   1828     0   /bin/bash
1.533   svstat              8787   8786     0   /command/svstat /service/nflx-httpd
1.533   perl                8788   8786     0   /usr/bin/perl -e $1=<>;$1=~/(\\d+) sec;/p...
[...]
```

# CPUs: runqlat

## Scheduler latency (run queue latency)

```
# runqlat.py 10 1
Tracing run queue latency... Hit Ctrl-C to end.
```

usecs	:	count	distribution
0 -> 1	:	1906	***
2 -> 3	:	22087	*****
4 -> 7	:	21245	*****
8 -> 15	:	7333	*****
16 -> 31	:	4902	*****
32 -> 63	:	6002	*****
64 -> 127	:	7370	*****
128 -> 255	:	13001	*****
256 -> 511	:	4823	*****
512 -> 1023	:	1519	**
1024 -> 2047	:	3682	*****
2048 -> 4095	:	3170	*****
4096 -> 8191	:	5759	*****
8192 -> 16383	:	14549	*****
16384 -> 32767	:	5589	*****

# CPUs: runqlen

## Run queue length

```
# runqlen.py 10 1
Sampling run queue length... Hit Ctrl-C to end.
```

runqlen	: count	distribution
0	: 47284	*****
1	: 211	
2	: 28	
3	: 6	
4	: 4	
5	: 1	
6	: 1	

# Memory: ffaults (book)

## Page faults by filename

```
# ffaults.bt
Attaching 1 probe...
^C

[...]
@[dpkg]: 18
@[sudoers.so]: 19
@[ld.so.cache]: 27
@[libpthread-2.27.so]: 29
@[ld-2.27.so]: 32
@[locale-archive]: 34
@[system.journal]: 39
@[libstdc++.so.6.0.25]: 43
@[libapt-pkg.so.5.0.2]: 47
@[BrowserMetrics-5D8A6422-77F1.pma]: 86
@[libc-2.27.so]: 168
@[i915]: 409
@[pkgcache.bin]: 860
@[]: 25038
```



# Disks: biolateny

## Disk I/O latency histograms, per second

```
# biolateny.py -mT 1 5
Tracing block device I/O... Hit Ctrl-C to end.
```

06:20:16

msecs	:	count	distribution
0 -> 1	:	36	*****
2 -> 3	:	1	*
4 -> 7	:	3	***
8 -> 15	:	17	*****
16 -> 31	:	33	*****
32 -> 63	:	7	*****
64 -> 127	:	6	*****

06:20:17

msecs	:	count	distribution
0 -> 1	:	96	*****
2 -> 3	:	25	*****
4 -> 7	:	29	*****

[...]

# File Systems: xfsslower

XFS I/O slower than a threshold (variants for ext4, btrfs, zfs)

```
# xfsslower.py 50
```

```
Tracing XFS operations slower than 50 ms
```

TIME	COMM	PID	T	BYTES	OFF_KB	LAT(ms)	FILENAME
21:20:46	java	112789	R	8012	13925	60.16	file.out
21:20:47	java	112789	R	3571	4268	136.60	file.out
21:20:49	java	112789	R	5152	1780	63.88	file.out
21:20:52	java	112789	R	5214	12434	108.47	file.out
21:20:52	java	112789	R	7465	19379	58.09	file.out
21:20:54	java	112789	R	5326	12311	89.14	file.out
21:20:55	java	112789	R	4336	3051	67.89	file.out
[...]							
22:02:39	java	112789	R	65536	1486748	182.10	shuffle_6_646_0.data
22:02:39	java	112789	R	65536	872492	30.10	shuffle_6_646_0.data
22:02:39	java	112789	R	65536	1113896	309.52	shuffle_6_646_0.data
22:02:39	java	112789	R	65536	1481020	400.31	shuffle_6_646_0.data
22:02:39	java	112789	R	65536	1415232	324.92	shuffle_6_646_0.data
22:02:39	java	112789	R	65536	1147912	119.37	shuffle_6_646_0.data
[...]							

# File Systems: xfsdist

## XFS I/O latency histograms by operation

```
# xfsdist.py 60
Tracing XFS operation latency... Hit Ctrl-C to end.

22:41:24:

operation = 'read'
      usecs      : count      distribution
      0 -> 1      : 382130    | ***** |
      2 -> 3      : 85717     | ***** |
      4 -> 7      : 23639     | **      |
      8 -> 15     : 5668      |          |
     16 -> 31     : 3594      |          |
     32 -> 63     : 21387     | **      |
[...]
```

operation	usecs	count	distribution
read	0 -> 1	382130	*****
	2 -> 3	85717	*****
	4 -> 7	23639	**
	8 -> 15	5668	
	16 -> 31	3594	
	32 -> 63	21387	**
write	0 -> 1	12925	*****
	2 -> 3	83375	*****

```
[...]
```

# Networking: tcplife

## TCP session lifespans with connection details

```
# tcplife.py
PID    COMM      LADDR      LPORT  RADDR      RPORT  TX_KB  RX_KB  MS
22597  recordProg 127.0.0.1   46644  127.0.0.1   28527   0      0  0.23
3277   redis-serv 127.0.0.1   28527  127.0.0.1   46644   0      0  0.28
22598  curl       100.66.3.172 61620  52.205.89.26 80      0      1  91.79
22604  curl       100.66.3.172 44400  52.204.43.121 80      0      1  121.38
22624  recordProg 127.0.0.1   46648  127.0.0.1   28527   0      0  0.22
3277   redis-serv 127.0.0.1   28527  127.0.0.1   46648   0      0  0.27
22647  recordProg 127.0.0.1   46650  127.0.0.1   28527   0      0  0.21
3277   redis-serv 127.0.0.1   28527  127.0.0.1   46650   0      0  0.26
[...]
```

# Networking: tcpsynbl (book)

# TCP SYN backlogs as histograms

```
# tcpsynbl.bt
Attaching 4 probes...
Tracing SYN backlog size. Ctrl-C to end.
^C
@backlog[backlog limit]: histogram of backlog size

@backlog[128]:
[0]                2 | @@@@

@backlog[500]:
[0]               2783 | @@@@
[1]                  9 |
[2, 4)              4 |
[4, 8)              1 |
```

# Languages: funccount

Count native function calls (C, C++, Go, etc)

```
# funccount.py 'tcp_s*'
Tracing 50 functions for "tcp_s*"... Hit Ctrl-C to end.
^C
FUNC                                COUNT
[...]
tcp_setsockopt                      1839
tcp_shutdown                        2690
tcp_sndbuf_expand                   2862
tcp_send_delayed_ack                9457
tcp_set_state                       10425
tcp_sync_mss                        12529
tcp_sendmsg_locked                  41012
tcp_sendmsg                         41236
tcp_send_mss                        42686
tcp_small_queue_check.isra.29      45724
tcp_schedule_loss_probe             64067
tcp_send_ack                        66945
tcp_stream_memory_free              178616
Detaching...
```

# Applications: mysqld\_qslower

## MySQL queries slower than a threshold

```
# mysqld_qslower.py $(pgrep mysqld)
Tracing MySQL server queries for PID 9908 slower than 1 ms...
TIME(s)      PID      MS QUERY
0.000000    9962    169.032 SELECT * FROM words WHERE word REGEXP '^bre.*n$'
1.962227    9962    205.787 SELECT * FROM words WHERE word REGEXP '^bpf.tools$'
9.043242    9962     95.276 SELECT COUNT(*) FROM words
23.723025    9962    186.680 SELECT count(*) AS count FROM words WHERE word REGEXP
'^bre.*n$'
30.343233    9962    181.494 SELECT * FROM words WHERE word REGEXP '^bre.*n$' ORDER BY
word
[...]
```

# Coping with so many BPF tools

- On all Netflix servers, **/apps/nflx-bpf-alltools** has tools from:
  - BCC, bpftrace, my book, Netflix internal
- Latest tools are fetched & installed in a hierarchy: cpu, disk, ...

```
bgregg@lgud-bgregg:~> ls --color ~/Git/nflx-bpf-alltools/root/apps/nflx-bpf-alltools/  
applications/      disk/              funcslower.py*    stackcount_example.txt  
argdist_example.txt  filesystems/      hypervisors/      stackcount.py*  
argdist.py*        funccount_example.txt  kernel/          tplist_example.txt  
bpflist_example.txt  funccount.py*     languages/        tplist.py*  
bpflist.py*        funclatency_example.txt  memory/         trace_example.txt  
containers/        funcslower.py*     networking/       trace.py*  
cpu/               funcslower_example.txt  security/
```

- An employee can look in “disk” for all disk tools.
- We are also building **GUIs** to front these tools

We might open source the nflx-bpf-alltools framework at some point



# **Performance Analysis with BPF**

# Start With Basics

- From my BCC tutorial:

<https://github.com/iovisor/bcc/blob/master/docs/tutorial.md>

1. uptime

2. dmesg | tail

3. vmstat 1

4. mpstat -P ALL 1

5. pidstat 1

6. iostat -xz 1

7. free -m

8. sar -n DEV 1

9. sar -n TCP,ETCP 1


10. top

# BCC Checklist

- ...continuing the BCC tutorial:
  1. execsnoop
  2. opensnoop
  3. ext4slower  
(or btrfs\*, xfs\*, zfs\*)
  4. biolatency
  5. biosnoop
  6. cachestat
  7. tcpconnect
  8. tcpaccept
  9. tcpretrans
  10. runqlat
  11. profile

# The **most important** skill to learn

- It's not the tools or the languages (bpftrace, BCC)
  - Analogy: good golf clubs don't make you a great golfer.
- It's the **know how**. It's years of experience and wisdom gained from practicing performance engineering.
- I've taught tracing to >1000 students as a professional instructor. What I found most effective was develop and teach performance analysis **methodologies**.
  - I documented many in my last book. Look out for 2<sup>nd</sup> ed.





# Methodologies

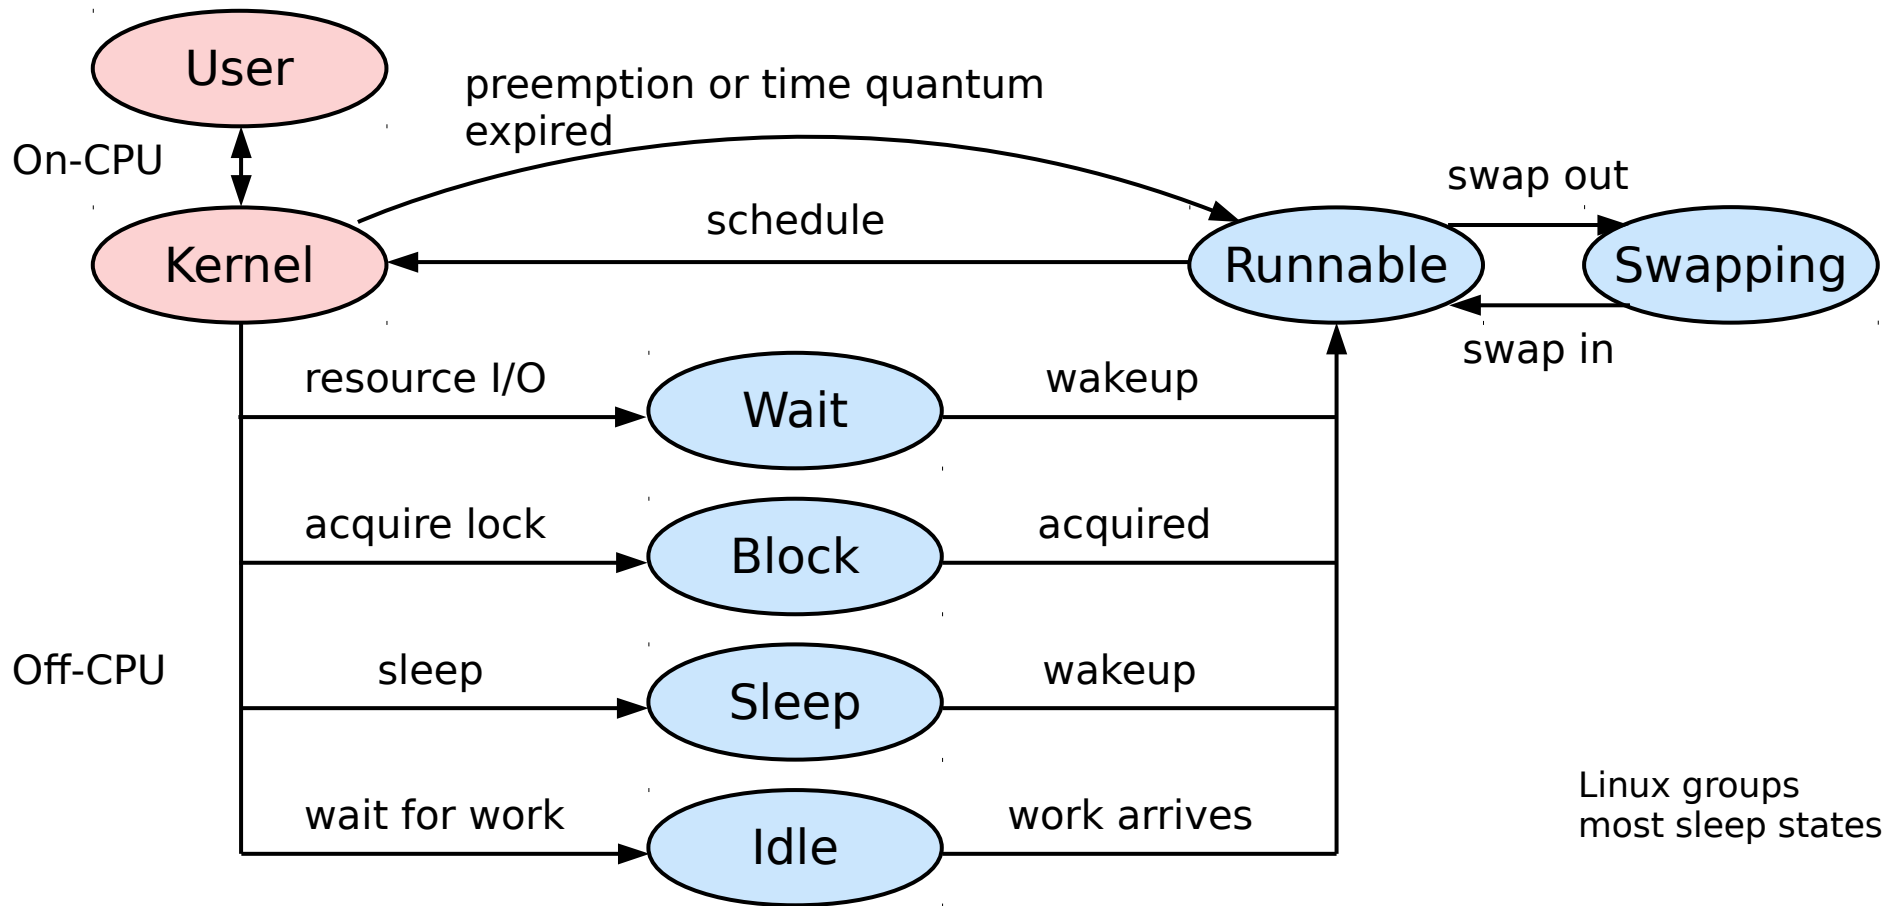
- Checklists
- Thread state analysis
- Workload characterization
- Reverse diagnosis
- Drill-down analysis
- Process of elimination
- 5 Whys
- (etc)

# Checklists

BPF Perf Tools:  
my diagram can  
be a checklist



# Thread State Analysis




# Workload Characterization

Analyze workload characteristics, not resulting performance

For example, CPUs:

1. **Who**: which PIDs, programs, users
2. **Why**: code paths, context
3. **What**: CPU instructions, cycles
4. **How**: changing over time





# Reverse Diagnosis

- Performance problems can have multiple causes and solutions. List them then find the metrics.
- Example: disk I/O bounded workload:
  - A) buy faster disks
  - B) change filesystem tuning to reduce I/O
  - C) change application config to improve I/O (e.g., more I/O threads)
  - D) change application to do I/O later (e.g., sync→async)
  - E) change application to cache I/O
  - F) change application to eliminate I/O (e.g., logic change)
- What metrics would identify each of these solutions?

# Other Methodologies

- Drill down analysis
- Process of elimination
- 5 Whys

# Lab 2

- `bpf-perf-workshop/lab2.md`

# Lab 2 Discussion

# Lab 3

- `bpf-perf-workshop/lab3.md`

# Lab 3 Discussion

**bpftrace**

# bpftrace Syntax

**bpftrace -e** 'k:do\_nanosleep /pid > 100/ { @[comm]++ }'

Probe

Filter  
(optional)

Action



# Probe Type Shortcuts

tracepoint	t	Kernel static tracepoints
usdt	U	User-level statically defined tracing
kprobe	k	Kernel function tracing
kretprobe	kr	Kernel function returns
uprobe	u	User-level function tracing
uretprobe	ur	User-level function returns
profile	p	Timed sampling across all CPUs
interval	i	Interval output
software	s	Kernel software events
hardware	h	Processor hardware events

# Filters

- `/pid == 181/`
- `/comm != "sshd"/`
- `/@ts[tid]/`

# Actions

- Per-event output
  - `printf()`
  - `system()`
  - `join()`
  - `time()`
- Map Summaries
  - `@ = count()` or `@++`
  - `@ = hist()`
  - ...

The following is in the [https://github.com/iovisor/bpftrace/blob/master/docs/reference\\_guide.md](https://github.com/iovisor/bpftrace/blob/master/docs/reference_guide.md)

# Functions

- **hist(n)** Log2 histogram
- **lhist(n, min, max, step)** Linear hist.
- **count()** Count events
- **sum(n)** Sum value
- **min(n)** Minimum value
- **max(n)** Maximum value
- **avg(n)** Average value
- **stats(n)** Statistics
- **str(s)** String
- **ksym(p)** Resolve kernel addr
- **usym(p)** Resolve user addr
- **kaddr(n)** Resolve kernel symbol
- **uaddr(n)** Resolve user symbol
- **printf(fmt, ...)** Print formatted
- **print(@x[, top[, div]])** Print map
- **delete(@x)** Delete map element
- **clear(@x)** Delete all keys/values
- **reg(n)** Register lookup
- **join(a)** Join string array
- **time(fmt)** Print formatted time
- **system(fmt)** Run shell command
- **cat(file)** Print file contents
- **exit()** Quit bpftrace

# Variable Types

- Basic Variables
  - `@global`
  - `@thread_local[tid]`
  - `$scratch`
- Associative Arrays
  - `@array[key] = value`
- Buitins
  - `pid`
  - `...`

# Builtin Variables

- **pid** Process ID (kernel tgid)
- **tid** Thread ID (kernel pid)
- **cgroup** Current Cgroup ID
- **uid** User ID
- **gid** Group ID
- **nsecs** Nanosecond timestamp
- **cpu** Processor ID
- **comm** Process name
- **kstack** Kernel stack trace
- **ustack** User stack trace
- **arg0, arg1, ...** Function args
- **retval** Return value
- **args** Tracepoint args
- **func** Function name
- **probe** Full probe name
- **curtask** Curr task\_struct (u64)
- **rand** Random number (u32)

# bpftrace: BPF observability front-end

## # Files opened by process

```
bpftrace -e 't:syscalls:sys_enter_open { printf("%s %s\n", comm,  
    str(args->filename)) }'
```

## # Read size distribution by process

```
bpftrace -e 't:syscalls:sys_exit_read { @[comm] = hist(args->ret) }'
```

## # Count VFS calls

```
bpftrace -e 'kprobe:vfs_* { @[func]++ }'
```

## # Show vfs\_read latency as a histogram

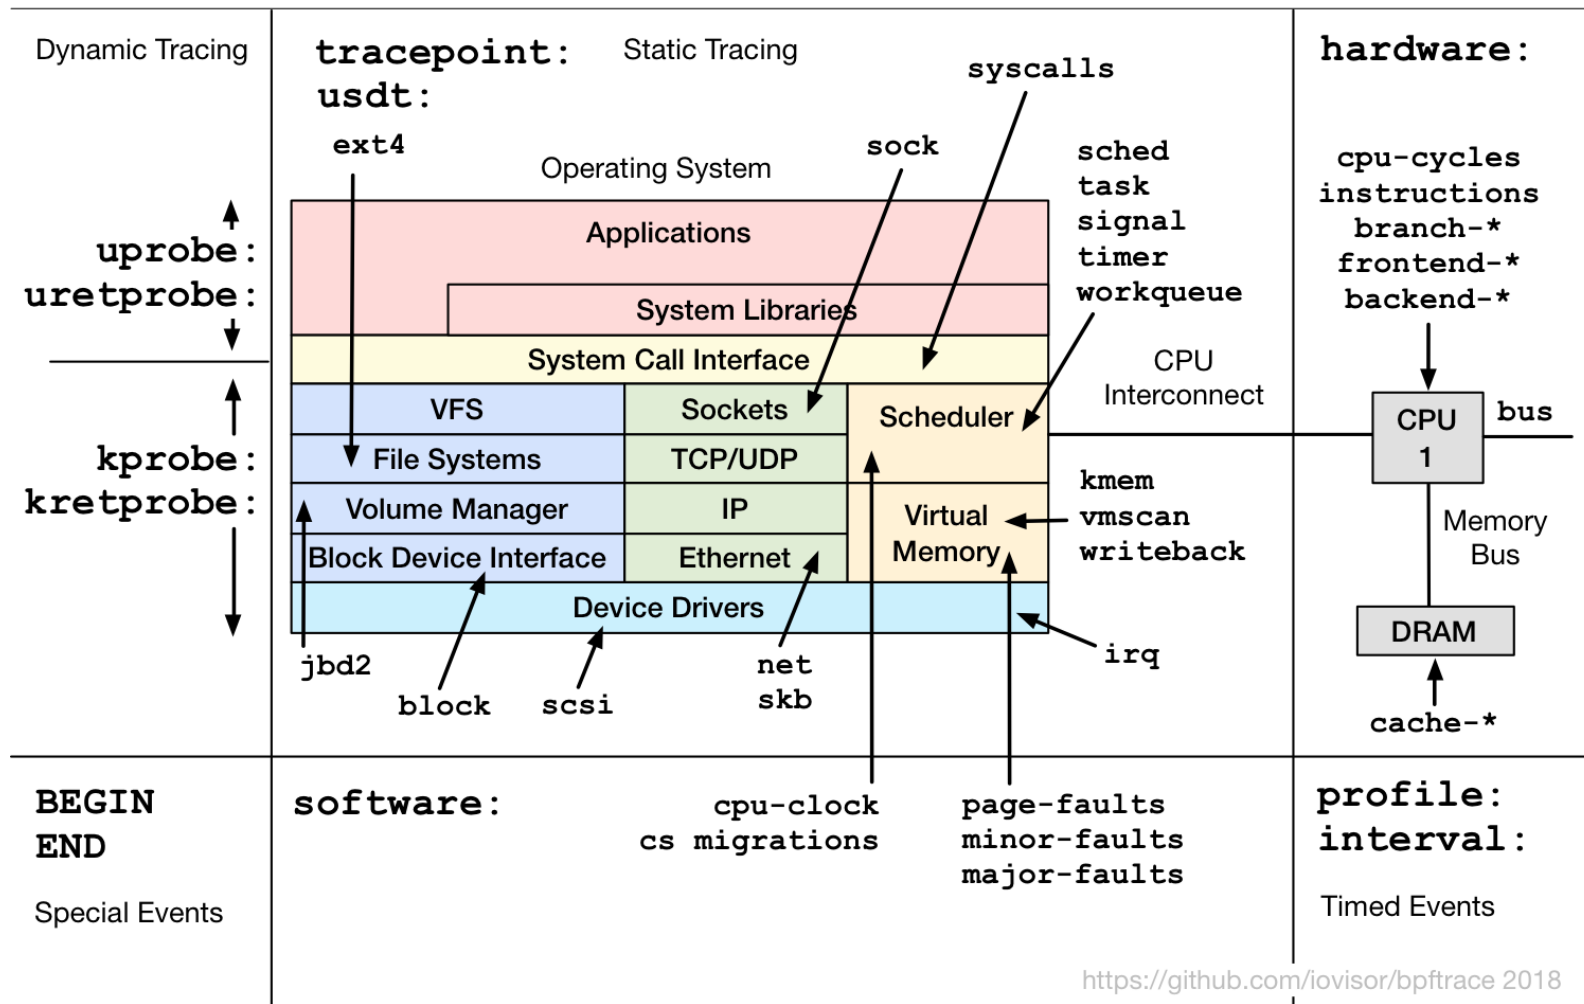
```
bpftrace -e 'k:vfs_read { @[tid] = nsecs }  
    kr:vfs_read /@[tid]/ { @ns = hist(nsecs - @[tid]); delete(@tid) }'
```

## # Trace user-level function

```
bpftrace -e 'uretprobe:bash:readline { printf("%s\n", str(retval)) }'
```

...

# Probes





# biolateny

```
#!/usr/local/bin/bpftrace

BEGIN
{
    printf("Tracing block device I/O... Hit Ctrl-C to end.\n");
}

kprobe:blk_account_io_start
{
    @start[arg0] = nsecs;
}

kprobe:blk_account_io_completion
/@start[arg0]/

{
    @usecs = hist((nsecs - @start[arg0]) / 1000);
    delete(@start[arg0]);
}
```

# >120 code examples

```
bpfttrace/tools> ls *.bt
```

bashreadline.bt	dc Snoop.bt	oomkill.bt	swapin.bt	tcpretrans.bt
biolateney.bt	execsnoop.bt	opensnoop.bt	syncsnoop.bt	tcpsynbl.bt
biosnoop.bt	gethostlatency.bt	pidpersec.bt	syscount.bt	threadsnoop.bt
biostacks.bt	killsnoop.bt	runqlat.bt	tcpaccept.bt	vfscount.bt
bitesize.bt	loads.bt	runqlen.bt	tcpconnect.bt	vfsstat.bt
capable.bt	mdflush.bt	setuids.bt	tcpdrop.bt	writeback.bt
cpuwalk.bt	naptime.bt	statsnoop.bt	tcplife.bt	xfsdist.bt

```
bpf-perf-tools-book/originals> ls */*.bt
```

Ch06_CPUs/cpufreq.bt	Ch10_Networking/skblife.bt
Ch06_CPUs/execsnoop.bt	Ch10_Networking/so1stbyte.bt
Ch06_CPUs/offcputime.bt	Ch10_Networking/soaccept.bt
Ch06_CPUs/runqlat.bt	Ch10_Networking/socketio.bt
Ch06_CPUs/runqlen.bt	Ch10_Networking/socksize.bt
Ch06_CPUs/smpcalls.bt	Ch10_Networking/sockstat.bt
Ch07_Memory/brkstack.bt	Ch10_Networking/soconnect.bt
Ch07_Memory/faults.bt	Ch10_Networking/soconnlat.bt
Ch07_Memory/ffaults.bt	Ch10_Networking/sofamily.bt
Ch07_Memory/hfaults.bt	Ch10_Networking/soprotocol.bt
[...]	

# Lab 4

- `bpf-perf-workshop/lab4.md`

Lab 5 is optional advanced exercises

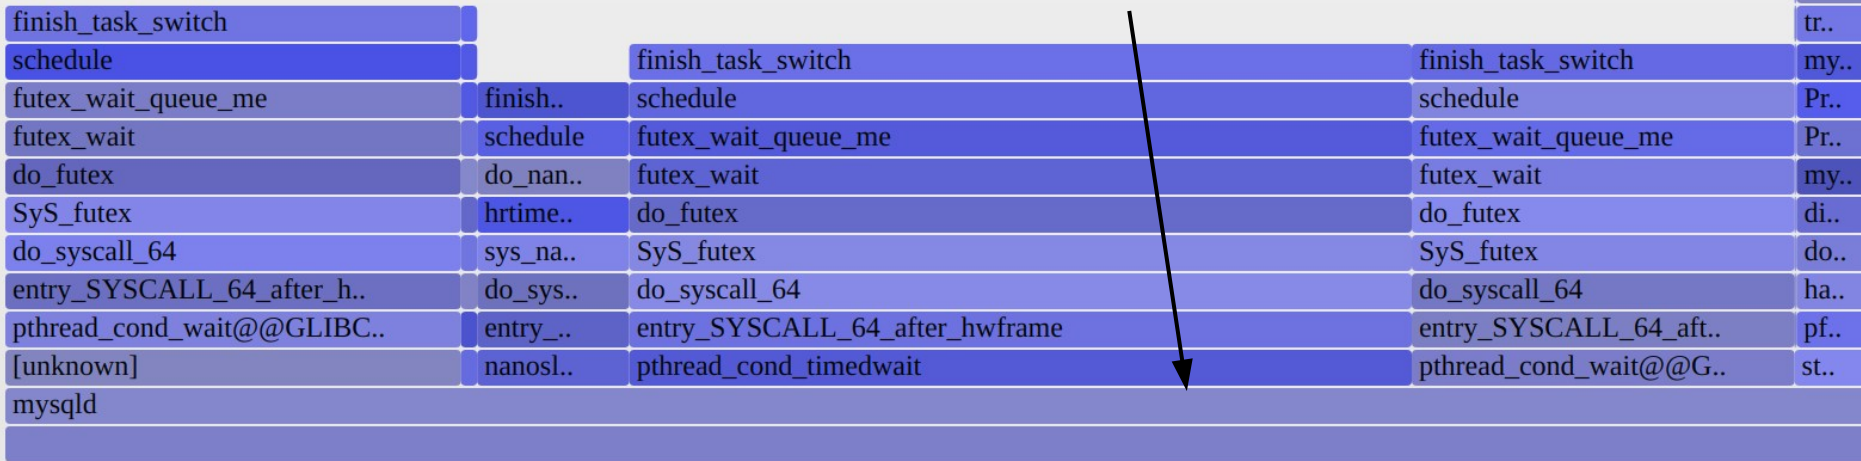
# Lab 4 & 5 Discussion

# Challenges

# Observability Challenges

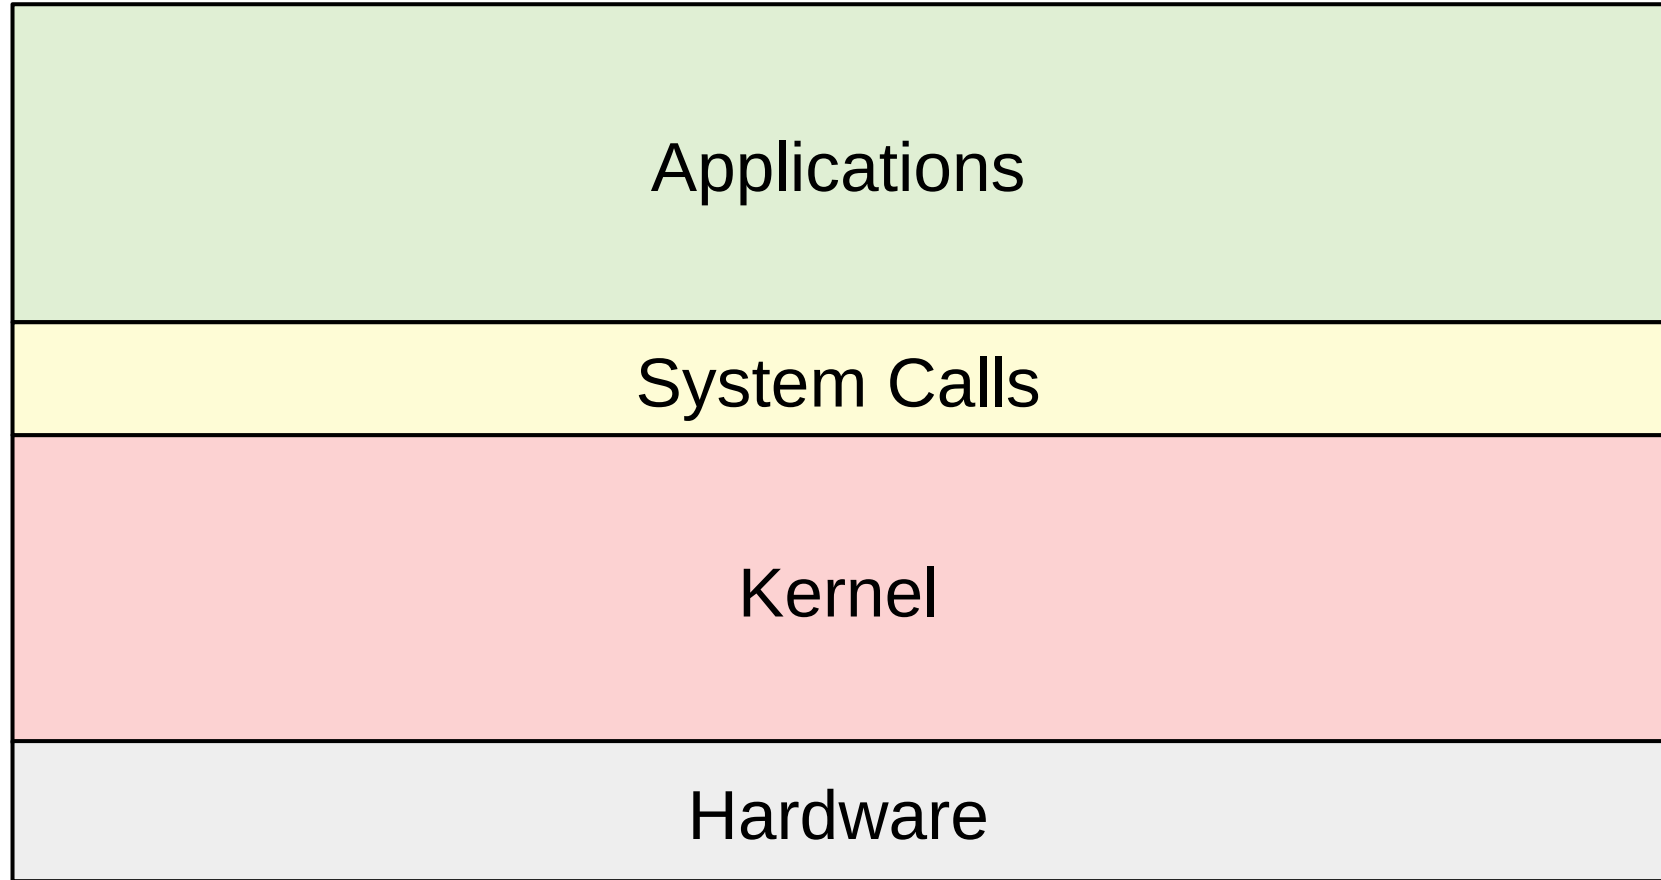
libc no frame pointer  
JIT function tracing

Broken off-CPU flame graph (no frame pointer)




**Future**

# 50 Years, one (dominant) OS model

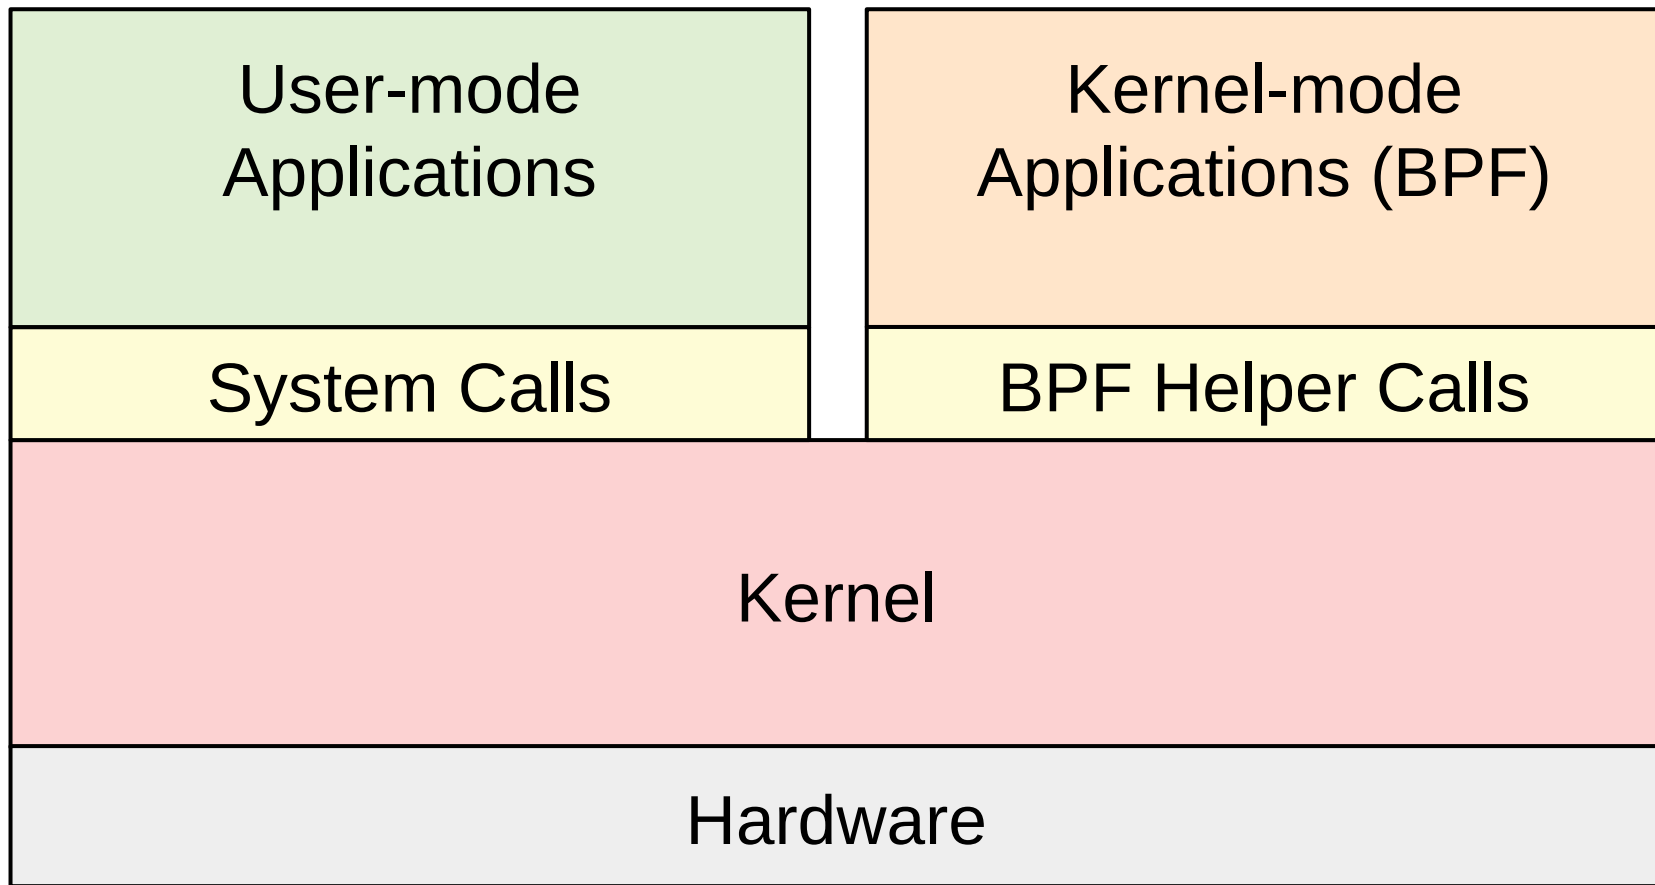




# Origins: Multics, 1960s



# Modern Linux: A new OS model



# Predictions

# Recap

# Learning Objectives

1. **slides:** Understand BPF, BCC, and bpftrace
2. **slides & discussion:** Follow different analysis methodologies
3. **lab1:** Use BCC tools to analyze disk I/O issues
4. **lab2:** “ “ short-lived process issues
5. **lab3:** “ “ runq latency issues
6. **lab4-5:** Develop at least one new bpftrace tool

# Thanks



BPF: Alexei Starovoitov, Daniel Borkmann, David S. Miller, Linus Torvalds, BPF community

BCC: Brenden Blanco, Yonghong Song, Sasha Goldsthein (who also does great workshops), BCC community

bpftime: Alastair Robertson, Matheus Marchini, Dan Xu, bpftime community

And thanks to Jérôme Petazzoni for tips on tutorials:

<http://jpetazzo.github.io/2015/09/10/how-to-deliver-great-tech-tutorials/>



USENIX LISA 2019, Portland, Oct 28-30